

Goal

We want to see how you would setup a small ELT pipeline, what you deem (un)important in such a setup.

Description of Use Case

Data have been extracted from an external data source and saved into a database as several tables. You will first need to extract the relevant data from these tables and perform some transformations to generate a single table with the properties defined below (normally this table would be used for ML tasks at paretos). You should use dbt to transform the data from the raw version into the final version.

Tasks

- Setup a git repository for this coding challenge.
- Setup a DB of your choice that can be used by dbt (PostgreSQL is recommended) and import the attached data
- Setup dbt to transform the imported tables:
 - Orders meta information
 - Order line items
 - Cost per SKU (Stock Keeping Unit) (some costs are missing)

Given the three tables above, think of a way to combine the information from multiple sources into a final (long or wide) table for daily reporting. It is up to you to decide what needs to be included in this table and how you organize your models. The two KPIs are very important for the end users of the table are:

- the margin per day: revenue - total cost of the SKU **in that order**
- the total number of orders per day: in order to better organize logistics

Note: Is there anything to report about the data? Inconsistencies?

- Is there anything else that you would have liked to do but did not have the time for?

Deliverables

- The git repo with the code+docs you developed. Zip it and send it to us via mail or invite us to the (private) repo.

Presentation

We'll schedule a call (30 min presentation, 10 min questions) where you can present your setup. During the presentation focus on the WHY, not the WHAT.

Additional information

- <https://www.getdbt.com/>